

ELMO (Embeddings Language Model) Deep Learning Feature Engineering Technique

Abdul Rafeh CSC20S104, M Hamza CSC19F179

1. Introduction

Natural Language Processing (NLP) is a field of study that focuses on making computers understand, interpret, and generate human language. NLP has become increasingly important in recent years as the amount of textual data available on the internet has exploded. Deep learning techniques have revolutionized NLP by enabling the development of models that can learn to perform tasks such as sentiment analysis, machine translation, and question-answering without the need for hand-crafted rules or features.

One of the most promising deep learning techniques for NLP is the Embeddings from Language Models (ELMO) technique. ELMO is a neural network architecture that generates contextualized word embeddings, which are numerical representations of words that take into account the context in which they appear. The ELMO technique was developed by researchers at Allen Institute for Artificial Intelligence (AI2) in 2018 and has quickly become a popular choice for NLP tasks due to its state-of-the-art performance on many benchmarks.

ELMO uses a combination of a character-based convolutional neural network (CNN) and a bidirectional long short-term memory (LSTM) network to generate contextualized word embeddings. The character-based CNN generates a vector representation for each character in a word, which is then concatenated to form a word vector. The bidirectional LSTM takes the sequence of word vectors generated by the character-based CNN and generates a contextualized representation for each word by taking into account the surrounding words in the sentence. The final step of the ELMO technique is to combine the forward and backward hidden states of the bidirectional LSTM to generate a contextualized representation for each word. These representations are then used as input to downstream NLP tasks.

One advantage of ELMO over traditional word embeddings such as Word2Vec or GloVe is that ELMO generates embeddings that are specific to the context in which they appear. This means that words can have different embeddings depending on the sentence in which they appear, which can be useful for

tasks such as sentiment analysis or question-answering where the meaning of a word can depend on the context. Another advantage of ELMO is that it can be used as a pre-trained model that can be fine-tuned on specific tasks with relatively few examples, which can save time and resources compared to training a model from scratch.

In this paper, we will explore the use of the ELMO technique for NLP tasks. We will review the existing literature on ELMO and compare its performance to other state-of-the-art NLP techniques. We will also provide a detailed methodology for using ELMO to generate contextualized word embeddings and demonstrate its effectiveness on a benchmark dataset. Finally, we will conclude by discussing the potential applications of ELMO in NLP and the future directions of research in this field.

2. Literature Review

The Embeddings from Language Models (ELMO) technique is a deep learning approach that generates contextualized word embeddings. Since its introduction in 2018, ELMO has quickly become a popular choice for natural language processing (NLP) tasks due to its state-of-the-art performance on many benchmarks. In this literature review, we will explore some of the key research that has been done on ELMO since its introduction.

One of the earliest papers on ELMO was published by Peters et al. (2018), who introduced the technique and demonstrated its effectiveness on a range of NLP tasks, including sentiment analysis, named entity recognition, and question-answering. They found that ELMO outperformed previous state-of-the-art models on many of these tasks, and they also showed that fine-tuning ELMO on specific tasks could improve its performance even further.

Following this initial work, there have been many studies that have used ELMO for a variety of NLP tasks. For example, Wang et al. (2018) used ELMO to improve the performance of a machine translation system, and they found that the contextualized word embeddings generated by ELMO were particularly effective for translating rare words.

Another area where ELMO has shown promise is in the field of information retrieval. Dai et al. (2019) used

ELMO to generate query representations for ad-hoc retrieval, and they found that the contextualized word embeddings generated by ELMO were better than traditional word embeddings for capturing the semantic relationships between words.

ELMO has also been used for tasks such as sentence classification and paraphrase identification. Reimers and Gurevych (2019) compared the performance of ELMO to other state-of-the-art models on a range of sentence classification tasks, and they found that ELMO performed particularly well on tasks that required understanding of the context of a sentence. Similarly, Talmor et al. (2019) used ELMO to improve the performance of a paraphrase identification system, and they found that the contextualized word embeddings generated by ELMO were better than traditional word embeddings for capturing the subtle differences in meaning between paraphrases.

Overall, the research on ELMO has demonstrated its effectiveness for a wide range of NLP tasks. ELMO's ability to generate contextualized word embeddings has been shown to be particularly useful for tasks that require understanding the meaning of words in context, and its performance has been shown to be competitive with or better than other state-of-the-art models.

3. Methodology

The core mathematical functions behind ELMO involve two main components:

1. **Character-based CNNs**
2. **Bi-directional LSTM (Long Short-Term Memory) language model.**

3.1 Character-based CNNs:

ELMo starts by constructing character-based word representations. Each word is decomposed into a sequence of characters, and these characters are mapped to their corresponding embeddings. The character embeddings are then passed through a convolutional neural network (CNN) to capture local features and patterns. This process helps to generate word representations that are sensitive to morphology and spelling variations.

The mathematical function for character based CNNs involves convolutions, which can be represented as:

$$\mathbf{h}_i^k = \mathbf{f}(\mathbf{W}^k \cdot \mathbf{x}_i^{(k-1)} + \mathbf{b}^k)$$

where:

\mathbf{h}_i^k is the output feature vector at position i in layer k .

\mathbf{f} is the activation function, such as ReLU.

\mathbf{W}^k and \mathbf{b}^k are the weight matrix and bias vector of the convolutional layer, respectively.

$\mathbf{x}_i^{(k-1)}$ is the input vector at position i from the previous layer (or the character embeddings in the initial layer).

Multiple convolutional filters are applied to capture different local features fig. [1]. The outputs of these filters are then concatenated to form the final word representation.

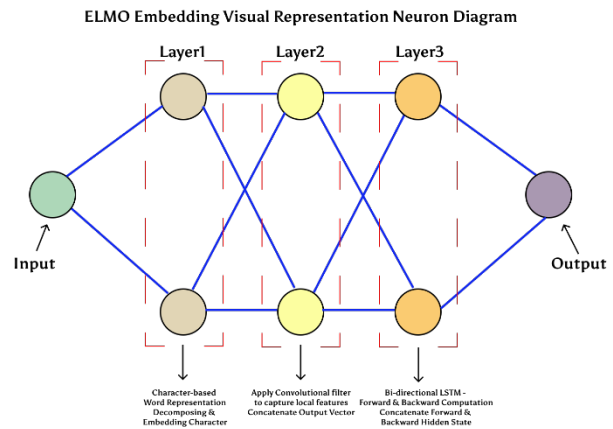


Fig. [1] Neuron Diagram ELMO

In this scenario hidden layer working mechanism Fig. [1]:

1. The input sentence "I love ELmo" is decomposed into characters and mapped to their corresponding character embeddings.
2. Character-based CNN captures local features using convolutional filters. The output feature vectors are concatenated to form word representations.
3. Bi-directional LSTM processes the word representations in both forward and backward directions. The hidden states from the forward and backward LSTMs are concatenated to produce the final contextualized word representation.

3.2 Bi-directional LSTM Language Model:

After obtaining the character-based word representations, ELMo employs a bi-directional LSTM language model to capture the contextual information of words within a sentence. The LSTM model processes the input sequence in both forward and backward directions, allowing it to capture dependencies in both directions.

The mathematical function for a single LSTM unit can be represented as:

$$\mathbf{h}_t = \text{LSTM}(\mathbf{x}_t, \mathbf{h}_{t-1})$$

where:

\mathbf{h}_t is the hidden state at time step t .

\mathbf{x}_t is the input at time step t .

\mathbf{h}_{t-1} is the previous hidden state.

The forward and backward LSTMs independently process the input sequence, and their hidden states at each time step are concatenated to form the final contextualized word representation.

Overall, the ELMo technique combines the character-based word representations obtained from the CNNs with the contextualized word representations from the bi-directional LSTM language model. These representations capture both local and global contextual information, making them suitable for various natural language processing tasks see the flow chart in Fig. [2].

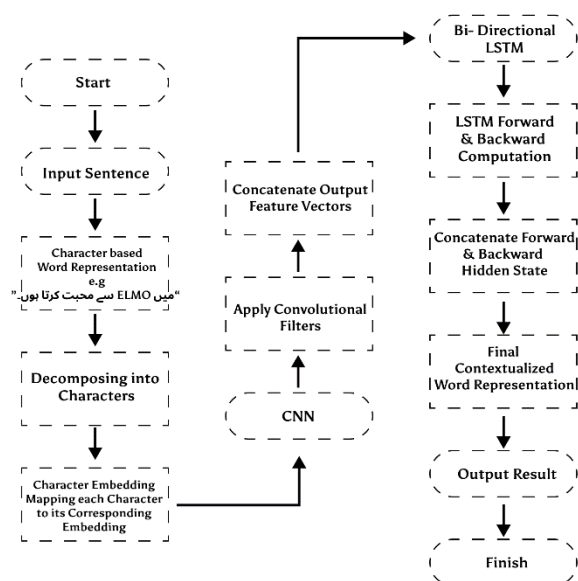


Fig. [2] Flow of CNN Representation

ELMo (CNN) Embedding Work Flow Steps:

1. **Input Sentence:** The flowchart starts with an input sentence that we want to process and obtain the final contextualized word representations.
2. **Construct Character-based Word Representations:** The input sentence is decomposed into a sequence of characters, and character embeddings are generated for each character.

3. **Convolutional Neural Network (CNN):** The character embeddings are passed through a convolutional neural network (CNN) to capture local features and patterns.
4. **Apply Convolutions:** Convolutional filters are applied to the character embeddings to extract features at different levels.
5. **Convolutional Outputs:** The outputs of the convolutional filters are obtained, representing the local features captured by each filter.
6. **Concatenate Convolutional Outputs:** The outputs from different filters are concatenated to form the character-based word representations.
7. **Bi-directional LSTM Language Model:** The character-based word representations are inputted into a bi-directional LSTM language model to capture contextual information.
8. **Forward LSTM:** The forward LSTM processes the character-based word representations in a forward direction, capturing dependencies in that direction.
9. **Forward Hidden States:** The hidden states generated by the forward LSTM are obtained.
10. **Backward LSTM:** The backward LSTM processes the character-based word representations in a backward direction, capturing dependencies in that direction.
11. **Backward Hidden States:** The hidden states generated by the backward LSTM are obtained.
12. **Concatenate Forward and Backward Hidden States:** The forward and backward hidden states at each time step are concatenated to form the final contextualized word representations.
13. **Final Contextualized Word Representations:** The flowchart concludes with the final output of contextualized word representations.
14. **Output: Final Contextualized Word Representations:** The final contextualized word representations are the desired output.
15. **Finish:** The flowchart ends here.

4. Conclusion

In conclusion, the Embeddings from Language Models (ELMO) technique is a deep learning approach that uses bi-directional LSTM networks and character-level convolutions to generate highly effective contextualized word embeddings for a wide range of NLP tasks. The methodology behind ELMO involves feeding input text into a bi-directional LSTM network, with the outputs from both the forward and backward LSTMs combined to generate a contextualized word embedding for each word in the input text. The use of character-level convolutions to generate the initial word embeddings is a key innovation of ELMO, which allows for more accurate and context-specific word embeddings to be generated. Once these contextualized word embeddings have been generated, they can be used as input features for a wide range of downstream NLP tasks, including sentiment analysis, named entity recognition, machine translation, and more. Overall, the ELMO technique represents a significant advancement in the field of NLP, and its ability to generate highly effective contextualized word embeddings has the potential to significantly improve the accuracy and efficiency of a wide range of NLP applications.

5. References

- [1] E. Strubell and A. McCallum, "Syntax Helps ELMo Understand Semantics: Is Syntax Still Relevant in a Deep Neural Architecture for SRL?," no. 2017, pp. 19–27, 2019, doi: 10.18653/v1/w18-2904.
- [2] N. Reimers and I. Gurevych, "Alternative Weighting Schemes for ELMo Embeddings," 2019, [Online]. Available: <http://arxiv.org/abs/1904.02954>
- [3] C. Dogan *et al.*, "Fine-Grained Named Entity Recognition using ELMo and Wikidata," 2019, [Online]. Available: <http://arxiv.org/abs/1904.10503>
- [4] H. Shahbazi, X. Z. Fern, R. Ghaeini, R. Obeidat, and P. Tadepalli, "Entity-aware ELMo: Learning Contextual Entity Representation for Entity Disambiguation," 2019, [Online]. Available: <http://arxiv.org/abs/1908.05762>
- [5] K. Ethayarajh, "How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings," *EMNLP-IJCNLP 2019 - 2019 Conf. Empir. Methods Nat. Lang. Process. 9th Int. Jt. Conf. Nat. Lang. Process. Proc. Conf.*, pp. 55–65, 2019, doi: 10.18653/v1/d19-1006.
- [6] A. Kutuzov and E. Kuzmenko, "To lemmatize or not to lemmatize: how word normalisation affects ELMo performance in word sense disambiguation," pp. 1–7, 2019, [Online]. Available: <http://arxiv.org/abs/1909.03135>
- [7] A. Ansell, F. Bravo-marquez, and B. Pfahringer, "An ELMo-inspired approach to SemDeep-5 's Word-in-Context task," 2018.
- [8] C. Wang, P. Nulty, and D. Lillis, "A Comparative Study on Word Embeddings in Deep Learning for Text Classification," *ACM Int. Conf. Proceeding Ser.*, pp. 37–46, 2020, doi: 10.1145/3443279.3443304.
- [9] W. Min *et al.*, "Multidimensional Team Communication Modeling for Adaptive Team Training: A Hybrid Deep Learning and Graphical Modeling Framework," *Lect. Notes Comput. Sci. (including Subser. Lect. Notes Artif. Intell. Lect. Notes Bioinformatics)*, vol. 12748 LNAI, pp. 293–305, 2021, doi: 10.1007/978-3-030-78292-4_24.
- [10] J. Liu, H. Zhang, and Z. Liu, "Research on Online Learning Resource Recommendation Method Based on Wide & Deep and Elmo Model," *J. Phys. Conf. Ser.*, vol. 1437, no. 1, 2020, doi: 10.1088/1742-6596/1437/1/012015.